

Financial Data Mining

Homework 3

Yen-Hsiu Chang

4.

(a)

```
# (i)
rm(list=ls())
N <- 1000
Pi <- c(.25,.5)
Delta <- c(0,1,2,3)
P <- c(1,3,5,10)
x.train.list <- list()
x.test.list <- list()
coef.list <- list()
train.error.logistic <- matrix(NA,1000,32)
test.error.logistic <- matrix(NA,1000,32)
train.error.lda <- matrix(NA,1000,32)
test.error.lda <- matrix(NA,1000,32)
set.seed(5)
library(glmnet)
y.train.1 <- c(rep(0,300*(1-Pi[1])),rep(1,300*Pi[1]))
y.train.2 <- c(rep(0,300*(1-Pi[2])),rep(1,300*Pi[2]))
y.test.1 <- c(rep(0,1000*(1-Pi[1])),rep(1,1000*Pi[1]))
y.test.2 <- c(rep(0,1000*(1-Pi[2])),rep(1,1000*Pi[2]))
for(s in seq(N)){
  for(pi in Pi){
    for(delta in Delta){
      for(p in P){
        # mean for two classes
        theta1 <- (delta/2) %%% c(1,rep(0,p-1)) # G = 0
        theta2 <- (-delta/2) %%% c(1,rep(0,p-1)) # G = 1
        # the mean matrix
        mu.train <- rbind(matrix(rep(theta1,(1-pi)*300),nrow=(1-pi)*300,
byrow=TRUE),
                        matrix(rep(theta2,pi*300),nrow=pi*300,byrow=TRUE))
        mu.test <- rbind(matrix(rep(theta1,(1-pi)*1000),nrow=(1-pi)*1000,
byrow=TRUE),
                        matrix(rep(theta2,pi*1000),nrow=pi*1000,byrow=TRUE))
        Sigma <- diag(rep(1,p))
        # observed data
        x.train <- mu.train + matrix(rnorm(300*p),300) %%% Sigma
        x.test <- mu.test + matrix(rnorm(1000*p),1000) %%% Sigma
        x.train.list <- c(x.train.list, list(x.train))
        x.test.list <- c(x.test.list, list(x.test))
      }
    }
  }
}
```

```

    }
  }
  ## (ii) logistic regression
  # record the estimated coefficients and the misclassification rate
  for(i in 1:32){
    if (i < 17){
      y.train <- y.train.1 ; y.test <- y.test.1
    }else{
      y.train <- y.train.2 ; y.test <- y.test.2
    }
    logistic <- glmnet(cbind(1,x.train.list[[i]]),y.train,family=c("binomial"),standardize=F ,lambda=0)
    coef.temp <- coef(logistic)[-2]
    coef.list <- c(coef.list,list(coef.temp))
    # compute Link
    a.train <- cbind(1,x.train.list[[i]]) %*% coef.temp
    a.test <- cbind(1,x.test.list[[i]]) %*% coef.temp
    # posterior probability
    p.train <- exp(a.train)/(1+exp(a.train))
    p.test <- exp(a.test)/(1+exp(a.test))
    # training error
    train.err <- mean((p.train>.5)!=y.train); train.error.logistic[s,i] <- train.err
    # test error
    test.err <- mean((p.test>.5)!=y.test) ; test.error.logistic[s,i] <- test.err
  }

  ## (iii) LDA
  for(i in 1:32){
    if (i < 17){
      g <- y.train <- y.train.1 ; y.test <- y.test.1
    }else{
      g <- y.train <- y.train.2 ; y.test <- y.test.2
    }
    # estimate parameters
    n1 <- sum(g==1) ; n2 <- sum(g==0) ; n <- n1+n2
    pi1.hat <- mean(g==1)
    pi2.hat <- mean(g==0)
    if(dim(x.train.list[[i]])[2]==1){
      mu.1.hat <- mean(x.train.list[[i]][g==1])
      mu.2.hat <- mean(x.train.list[[i]][g==0])
      # within class covariance
      S.w <- (t(x.train.list[[i]][g==1] - rep(1,n1)%*% t(mu.1.hat)) %*%
        (x.train.list[[i]][g==1] - rep(1,n1)%*% t(mu.1.hat)) +
        t(x.train.list[[i]][g==0] - rep(1,n2)%*% t(mu.2.hat)) %*%
        (x.train.list[[i]][g==0] - rep(1,n2)%*% t(mu.2.hat)))/(n-2)
    }
  }

```

```

    }else{
      mu.1.hat <- apply(x.train.list[[i]][g==1,],2,mean)
      mu.2.hat <- apply(x.train.list[[i]][g==0,],2,mean)
      # within class covariance
      S.w <- (t(x.train.list[[i]][g==1,] - rep(1,n1)%*% t(mu.1.hat)) %
*%
          (x.train.list[[i]][g==1,] - rep(1,n1)%*% t(mu.1.hat)) +
*%
          t(x.train.list[[i]][g==0,] - rep(1,n2)%*% t(mu.2.hat)) %
          (x.train.list[[i]][g==0,] - rep(1,n2)%*% t(mu.2.hat)))/(
n-2)
    }
    # decision boundary
    coefficient <- solve(S.w) %*% (mu.2.hat-mu.1.hat)
    intercept <- log(pi2.hat/pi1.hat) - t(mu.2.hat+mu.1.hat) %*% solve
(S.w) %*%
      (mu.2.hat-mu.1.hat)/2
    # misclassification rate
    g.hat.train <- ((x.train.list[[i]] %*% coefficient + as.vector(in
tercept)) < 0) * 1
    train.err <- mean(g.hat.train != g) ; train.error.lda[s,i] <- trai
n.err
    g.hat.test <- ((x.test.list[[i]] %*% coefficient + as.vector(inte
rcept)) < 0) * 1
    test.err <- mean(g.hat.test != y.test) ; test.error.lda[s,i] <- te
st.err
  }
  # empty the lists
  x.train.list <- c()
  x.test.list <- c()
}

```

(b)

```

train.error.logistic.avg <- apply(train.error.logistic,2,mean)
test.error.logistic.avg <- apply(test.error.logistic,2,mean)
train.error.lda.avg <- apply(train.error.lda,2,mean)
test.error.lda.avg <- apply(test.error.lda,2,mean)
train.err.log.avg <- matrix(train.error.logistic.avg,nrow=4)
test.err.log.avg <- matrix(test.error.logistic.avg,nrow=4)
train.err.lda.avg <- matrix(train.error.lda.avg,nrow=4)
test.err.lda.avg <- matrix(test.error.lda.avg,nrow=4)

```



```

logistic <- glmnet(cbind(1,x.train),y.train,family=c("binomial"),lambda
=0)
coef.temp <- coef(logistic)[-2]
# compute link
a.train <- cbind(1,x.train) %*% coef.temp
a.test <- cbind(1,x.test) %*% coef.temp
# posterior probability
p.train <- exp(a.train)/(1+exp(a.train))
p.test <- exp(a.test)/(1+exp(a.test))
# training error
train.err <- mean((p.train>.5)!=y.train)
train.err.logist <- c(train.err.logist,train.err)
# test error
test.err <- mean((p.test>.5)!=y.test)
test.err.logist <- c(test.err.logist,test.err)

# LDA
g <- y.train
n1 <- sum(g==1) ; n2 <- sum(g==0) ; n <- n1+n2
pi1.hat <- mean(g==1)
pi2.hat <- mean(g==0)
mu.1.hat <- mean(x.train[g==1])
mu.2.hat <- mean(x.train[g==0])
# within class covariance
S.w <- (t(x.train[g==1] - rep(1,n1)%*% t(mu.1.hat)) %*%
(x.train[g==1] - rep(1,n1)%*% t(mu.1.hat)) +
t(x.train[g==0] - rep(1,n2)%*% t(mu.2.hat)) %*%
(x.train[g==0] - rep(1,n2)%*% t(mu.2.hat)))/(n-2)
# decision boundary
coefficient <- solve(S.w) %*% (mu.2.hat-mu.1.hat)
intercept <- log(pi2.hat/pi1.hat) - t(mu.2.hat+mu.1.hat) %*% solve(S.w)
%*%
(mu.2.hat-mu.1.hat)/2
# misclassification rate
g.hat.train <- ((x.train %*% coefficient + as.vector(intercept)) < 0)
* 1
train.err <- mean(g.hat.train != g)
train.err.lda <- c(train.err.lda,train.err)
g.hat.test <- ((x.test %*% coefficient + as.vector(intercept)) < 0) *
1
test.err <- mean(g.hat.test != y.test)
test.err.lda <- c(test.err.lda,test.err)
}
mean(train.err.logist)

## [1] 0.09909667

mean(test.err.logist)

## [1] 0.100461

```

```
mean(train.err.lda)
```

```
## [1] 0.14372
```

```
mean(test.err.lda)
```

```
## [1] 0.144532
```

If the joint distribution of (X, G) is student t-distribution instead of normal distribution, logistic regression will work much better than LDA. That is because, for student t-distribution, observed data have more outliers, and these outliers will affect the estimation of common covariance matrix. This means LDA is not robust to gross outliers. From the simulation, we can find that misclassification rates of LDA are much higher than that of logistic regression.

5.


(a)

```
library(glmnet)
rm(list=ls())
load("hw3.RData")
x.train <- rbind(train2, train3, train8)
mu.hat <- apply(x.train, 2, mean)
x.centered <- x.train - rep(1, nrow(x.train))%*%t(mu.hat)
x.test <- rbind(test2, test3, test8)
x.test.centered <- x.test - rep(1, nrow(x.test))%*%t(mu.hat)
```

(b)

```
# PCA by SVD
x.svd <- svd(x.centered)
# PCA projection
fst.score.pca <- x.centered %*% x.svd$v[,1]
snd.score.pca <- x.centered %*% x.svd$v[,2]
n2 <- nrow(train2) ; n3 <- nrow(train3); n8 <- nrow(train8) ; n <- nrow(x.centered)

# PCA spectral
plot(x.svd$d)
```



(c)

```
# compute mean estimate: overall and group means
mu.hat <- apply(x.centered,2,mean)
mu.1.hat <- apply(x.centered[1:n2,],2,mean)
mu.2.hat <- apply(x.centered[(n2+1):(n2+n3),],2,mean)
mu.3.hat <- apply(x.centered[(n2+n3+1):n,],2,mean)

# between class covariance
S.b <- ((n2)*(mu.1.hat-mu.hat)%*%t(mu.1.hat-mu.hat)+
        (n3)*(mu.2.hat-mu.hat)%*%t(mu.2.hat-mu.hat)+
        (n8)*(mu.3.hat-mu.hat)%*%t(mu.3.hat-mu.hat))/(n-1)

# within class covariance
S.w <- (t(x.centered[1:n2,] - rep(1,n2)%*% t(mu.1.hat)) %*% (x.centered
[1:n2,] - rep(1,n2)%*% t(mu.1.hat)) +
        t(x.centered[(n2+1):(n2+n3),] - rep(1,n3)%*% t(mu.2.hat)) %*% (x
.centered[(n2+1):(n2+n3),] - rep(1,n3)%*% t(mu.2.hat)) +
        t(x.centered[(n2+n3+1):n,] - rep(1,n8)%*% t(mu.3.hat)) %*% (x
.centered[(n2+n3+1):n,] - rep(1,n8)%*% t(mu.3.hat)))/(n-3)

# total variance
S.t <- t(x.centered - rep(1,n)%*% t(mu.hat)) %*% (x.centered - rep(1,n)
%*% t(mu.hat))

# relation
# S.t - S.b*(n-1) - S.w*(n-K) = 0
max(abs(S.t - S.b*(n-1) - S.w*(n-3)))

## [1] 5.547918e-11
```

```


# define relative matrix
S <- solve(S.w) %*% S.b
# eigen decomp. of S
S.eig <- eigen(S)
t(Re(S.eig$vectors[,1:2]))%*%(Re(S.eig$vectors[,1:2])) # not orthogonal
any more

##           [,1]      [,2]
## [1,]  1.0000000000 -0.001460679
## [2,] -0.001460679  1.000000000

# retain the leading two scores
fst.score.fda <- x.centered %*% Re(S.eig$vectors[,1])
snd.score.fda <- x.centered %*% Re(S.eig$vectors[,2])

# FDA spectral
plot(Re(S.eig$values))

```



(d)

```

# OLS
g <- c(rep(1,n2),rep(2,n3),rep(3,n8))
N2 <- nrow(test2) ;N3 <- nrow(test3) ;N8 <- nrow(test8)
g.test <- c(rep(1,N2),rep(2,N3),rep(3,N8))
y.train <- cbind(g==1,g==2,g==3)*1
y.test <- cbind(g.test==1,g.test==2,g.test==3)*1
# check all the row sums are 1
# all(apply(y.train,1,sum)==1)

```



```

lm1 <- lm(y.train[,1]~.,data=data.frame(x.centered))
lm2 <- lm(y.train[,2]~.,data=data.frame(x.centered))
lm3 <- lm(y.train[,3]~.,data=data.frame(x.centered))

# prediction
y.train.hat <- cbind(fitted(lm1),fitted(lm2),fitted(lm3))
g.hat <- apply(y.train.hat,1,which.max)
# train error
ols.train.error <- mean(g.hat!=g)

y.test.hat <- cbind(predict(lm1,data.frame(x.test.centered)),predict(lm
2,data.frame(x.test.centered)),predict(lm3,data.frame(x.test.centered))
)
g.test.hat <- apply(y.test.hat,1,which.max)
# test error
ols.test.error <- mean(g.test.hat!=g.test)

# Logistic regression
ans.logistic <- glmnet(x.centered,g, family=c("multinomial"),lambda=0,s
tandardize=F)
g.hat.logist <- predict(ans.logistic,x.centered,type = "class")
# apply(predict(ans.logistic,x.centered,type = "response"),1,which.max)
g.test.hat.logist <- predict(ans.logistic,x.test.centered,type = "class
")

# train error
logistic.train.error <- mean(g.hat.logist!=g)
# test error
logistic.test.error <- mean(g.test.hat.logist!=g.test)

# LDA
pi1.hat <- mean(g==1)
pi2.hat <- mean(g==2)
pi3.hat <- mean(g==3)

mu.1.hat <- apply(x.centered[1:n2,],2,mean)
mu.2.hat <- apply(x.centered[(n2+1):(n2+n3),],2,mean)
mu.3.hat <- apply(x.centered[(n2+n3+1):n,],2,mean)

S.w <- (t(x.centered[1:n2,] - rep(1,n2)%*% t(mu.1.hat)) %*% (x.centered
[1:n2,] - rep(1,n2)%*% t(mu.1.hat)) +
      t(x.centered[(n2+1):(n2+n3),] - rep(1,n3)%*% t(mu.2.hat)) %*% (x
.centered[(n2+1):(n2+n3),] - rep(1,n3)%*% t(mu.2.hat)) +
      t(x.centered[(n2+n3+1):n,] - rep(1,n8)%*% t(mu.3.hat)) %*% (x.ce
ntered[(n2+n3+1):n,] - rep(1,n8)%*% t(mu.3.hat)))/(n-3)

slope1 <- solve(S.w) %*% (mu.2.hat-mu.1.hat)
intercept1 <- log(pi2.hat/pi1.hat) - t(mu.2.hat+mu.1.hat) %*% solve(S.w

```

```

) %*% (mu.2.hat-mu.1.hat)/2

slope2 <- solve(S.w) %*% (mu.3.hat-mu.1.hat)
intercept2 <- log(pi3.hat/pi1.hat) - t(mu.3.hat+mu.1.hat) %*% solve(S.w)
) %*% (mu.3.hat-mu.1.hat)/2

slope3 <- solve(S.w) %*% (mu.3.hat-mu.2.hat)
intercept3 <- log(pi3.hat/pi2.hat) - t(mu.3.hat+mu.2.hat) %*% solve(S.w)
) %*% (mu.3.hat-mu.2.hat)/2
# train error
delta.hat.train <-
cbind(log(pi1.hat)-diag((x.centered-rep(1,n)%*% t(mu.1.hat))%*% solve(S
.w)%*%t(x.centered-rep(1,n)%*% t(mu.1.hat))/2),
      log(pi2.hat)-diag((x.centered-rep(1,n)%*% t(mu.2.hat))%*% solve(S
.w)%*%t(x.centered-rep(1,n)%*% t(mu.2.hat))/2),
      log(pi3.hat)-diag((x.centered-rep(1,n)%*% t(mu.3.hat))%*% solve(S
.w)%*%t(x.centered-rep(1,n)%*% t(mu.3.hat))/2))
g.hat.lda <- apply(delta.hat.train,1,which.max)

lda.train.error <- mean(g.hat.lda!=g)

# test error
N <- N2+N3+N8
delta.hat.test <-
cbind(log(pi1.hat)-diag((x.test.centered-rep(1,N)%*% t(mu.1.hat))%*% so
lve(S.w)%*%t(x.test.centered-rep(1,N)%*% t(mu.1.hat))/2),
      log(pi2.hat)-diag((x.test.centered-rep(1,N)%*% t(mu.2.hat))%*% so
lve(S.w)%*%t(x.test.centered-rep(1,N)%*% t(mu.2.hat))/2),
      log(pi3.hat)-diag((x.test.centered-rep(1,N)%*% t(mu.3.hat))%*% so
lve(S.w)%*%t(x.test.centered-rep(1,N)%*% t(mu.3.hat))/2))
g.test.hat.lda <- apply(delta.hat.test,1,which.max)
lda.test.error <- mean(g.test.hat.lda!=g.test)

# error table
error.table <- data.frame(OLS=c(ols.train.error,ols.test.error),Logistic=
c(logistic.train.error,logistic.test.error),LDA=c(lda.train.error,lda
.test.error))
row.names(error.table)<-c("train error",'test error');error.table

##              OLS      Logistic      LDA
## train error 0.02278612 0.00000000 0.02382185
## test error  0.08490566 0.08301887 0.08867925

```

See the table.

(e)

```

x.train.pc <- cbind(fst.score.pca,snd.score.pca)
x.test.pc <- cbind(x.test.centered%*%x.svd$v[,1],x.test.centered%*%x.sv
d$v[,2])

```

```

# OLS
# g <- c(rep(1,n2),rep(2,n3),rep(3,n8))
# N2 <- nrow(test2) ;N3 <- nrow(test3) ;N8 <- nrow(test8)
# g.test <- c(rep(1,N2),rep(2,N3),rep(3,N8))
# y.train <- cbind(g==1,g==2,g==3)*1
# y.test <- cbind(g.test==1,g.test==2,g.test==3)*1
# check all the row sums are 1
# all(apply(y.train,1,sum)==1)
lm1 <- lm(y.train[,1]~.,data=data.frame(x.train.pc))
lm2 <- lm(y.train[,2]~.,data=data.frame(x.train.pc))
lm3 <- lm(y.train[,3]~.,data=data.frame(x.train.pc))

# prediction
y.train.hat <- cbind(fitted(lm1),fitted(lm2),fitted(lm3))
g.hat <- apply(y.train.hat,1,which.max)
# train error
ols.train.error <- mean(g.hat!=g)

y.test.hat <- cbind(predict(lm1,data.frame(x.test.pc)),predict(lm2,data
.frame(x.test.pc)),predict(lm3,data.frame(x.test.pc)))
g.test.hat <- apply(y.test.hat,1,which.max)
# test error
ols.test.error <- mean(g.test.hat!=g.test)

# Logistic regression
ans.logistic <- glmnet(x.train.pc,g, family=c("multinomial"),lambda=0,s
tandardize=F)
g.hat.logist <- predict(ans.logistic,x.train.pc,type = "class")
# apply(predict(ans.logistic,x.centered,type = "response"),1,which.max)
g.test.hat.logist <- predict(ans.logistic,x.test.pc,type = "class")

# train error
logistic.train.error <- mean(g.hat.logist!=g)
# test error
logistic.test.error <- mean(g.test.hat.logist!=g.test)

# LDA
pi1.hat <- mean(g==1)
pi2.hat <- mean(g==2)
pi3.hat <- mean(g==3)

mu.1.hat <- apply(x.train.pc[1:n2,],2,mean)
mu.2.hat <- apply(x.train.pc[(n2+1):(n2+n3),],2,mean)
mu.3.hat <- apply(x.train.pc[(n2+n3+1):n,],2,mean)

S.w <- (t(x.train.pc[1:n2,] - rep(1,n2)%*% t(mu.1.hat)) %*% (x.train.pc
[1:n2,] - rep(1,n2)%*% t(mu.1.hat)) +
        t(x.train.pc[(n2+1):(n2+n3),] - rep(1,n3)%*% t(mu.2.hat)) %*% (x
.train.pc[(n2+1):(n2+n3),] - rep(1,n3)%*% t(mu.2.hat)) +

```

```

      t(x.train.pc[(n2+n3+1):n,] - rep(1,n8)%*% t(mu.3.hat)) %*% (x.train.pc[(n2+n3+1):n,] - rep(1,n8)%*% t(mu.3.hat)))/(n-3)

slope1 <- solve(S.w) %*% (mu.2.hat-mu.1.hat)
intercept1 <- log(pi2.hat/pi1.hat) - t(mu.2.hat+mu.1.hat) %*% solve(S.w) %*% (mu.2.hat-mu.1.hat)/2

slope2 <- solve(S.w) %*% (mu.3.hat-mu.1.hat)
intercept2 <- log(pi3.hat/pi1.hat) - t(mu.3.hat+mu.1.hat) %*% solve(S.w) %*% (mu.3.hat-mu.1.hat)/2

slope3 <- solve(S.w) %*% (mu.3.hat-mu.2.hat)
intercept3 <- log(pi3.hat/pi2.hat) - t(mu.3.hat+mu.2.hat) %*% solve(S.w) %*% (mu.3.hat-mu.2.hat)/2
# train error
delta.hat.train <-
cbind(log(pi1.hat)-diag((x.train.pc-rep(1,n)%*% t(mu.1.hat))%*% solve(S.w)%*%t(x.train.pc-rep(1,n)%*% t(mu.1.hat))/2),
      log(pi2.hat)-diag((x.train.pc-rep(1,n)%*% t(mu.2.hat))%*% solve(S.w)%*%t(x.train.pc-rep(1,n)%*% t(mu.2.hat))/2),
      log(pi3.hat)-diag((x.train.pc-rep(1,n)%*% t(mu.3.hat))%*% solve(S.w)%*%t(x.train.pc-rep(1,n)%*% t(mu.3.hat))/2))
g.hat.lda <- apply(delta.hat.train,1,which.max)

lda.train.error <- mean(g.hat.lda!=g)

# test error
N <- N2+N3+N8
delta.hat.test <-
cbind(log(pi1.hat)-diag((x.test.pc-rep(1,N)%*% t(mu.1.hat))%*% solve(S.w)%*%t(x.test.pc-rep(1,N)%*% t(mu.1.hat))/2),
      log(pi2.hat)-diag((x.test.pc-rep(1,N)%*% t(mu.2.hat))%*% solve(S.w)%*%t(x.test.pc-rep(1,N)%*% t(mu.2.hat))/2),
      log(pi3.hat)-diag((x.test.pc-rep(1,N)%*% t(mu.3.hat))%*% solve(S.w)%*%t(x.test.pc-rep(1,N)%*% t(mu.3.hat))/2))
g.test.hat.lda <- apply(delta.hat.test,1,which.max)
lda.test.error <- mean(g.test.hat.lda!=g.test)

# error table
error.table2 <- data.frame("OLS_PCA"=c(ols.train.error,ols.test.error),
                           "Logistic_PCA"=c(logistic.train.error,logistic.test.error),"LDA_PCA"=c(lda.train.error,lda.test.error))
row.names(error.table2)<-c("train error",'test error')


# plot the first two PC scores and the decision boundaries
# OLS
x <- cbind(1,x.train.pc)
B <- solve(t(x)%*%x)%*%t(x)%*%y.train

```

```

par(mfrow=c(1,3))
plot(fst.score.pca,snd.score.pca,col=rep(2:4,c(n2,n3,n8)),pch=rep(2:4,c(
n2,n3,n8)),xlab="First Score",ylab="Second Score",main="OLS")
abline(b=(B[2,2]-B[2,1])/(B[3,1]-B[3,2]),a=(B[1,2]-B[1,1])/(B[3,1]-B[3,
2]))
abline(b=(B[2,3]-B[2,2])/(B[3,2]-B[3,3]),a=(B[1,3]-B[1,2])/(B[3,2]-B[3,
3]))
abline(b=(B[2,3]-B[2,1])/(B[3,1]-B[3,3]),a=(B[1,3]-B[1,1])/(B[3,1]-B[3,
3]))
# logistic
coef_logit <- coef(ans.logistic)
coef.log <- cbind(coef(ans.logistic)[[1]],coef(ans.logistic)[[2]],coef(
ans.logistic)[[3]])
plot(fst.score.pca,snd.score.pca,col=rep(2:4,c(n2,n3,n8)),pch=rep(2:4,c(
n2,n3,n8))
,xlab="First Score",ylab="Second Score",main="Logistic Regression"
)
abline(b=(coef.log[2,2]-coef.log[2,1])/(coef.log[3,1]-coef.log[3,2])
,a=(coef.log[1,2]-coef.log[1,1])/(coef.log[3,1]-coef.log[3,2]))
abline(b=(coef.log[2,3]-coef.log[2,2])/(coef.log[3,2]-coef.log[3,3])
,a=(coef.log[1,3]-coef.log[1,2])/(coef.log[3,2]-coef.log[3,3]))
abline(b=(coef.log[2,3]-coef.log[2,1])/(coef.log[3,1]-coef.log[3,3])
,a=(coef.log[1,3]-coef.log[1,1])/(coef.log[3,1]-coef.log[3,3]))
# LDA
plot(fst.score.pca,snd.score.pca,col=rep(2:4,c(n2,n3,n8)),pch=rep(2:4,c(
n2,n3,n8))
,xlab="First Score",ylab="Second Score",main="LDA")
abline(b=-slope1[1]/slope1[2],a=-intercept1/slope1[2])
abline(b=-slope2[1]/slope2[2],a=-intercept2/slope2[2])
abline(b=-slope3[1]/slope3[2],a=-intercept3/slope3[2])

```



error.table2

##	OLS_PCA	Logistic_PCA	LDA_PCA
## train error	0.1532885	0.1460383	0.1589850
## test error	0.2396226	0.2377358	0.2320755

See the table.

(f)

```
x.train.fda <- cbind(fst.score.fda,snd.score.fda)
x.test.fda <- cbind(x.test.centered%%Re(S.eig$vector[,1]),x.test.cent
ered%%Re(S.eig$vector[,2]))
# t(Re(S.eig$vector[,1:2]))%%(Re(S.eig$vector[,1:2]))
# OLS
# g <- c(rep(1,n2),rep(2,n3),rep(3,n8))
# N2 <- nrow(test2) ;N3 <- nrow(test3) ;N8 <- nrow(test8)
# g.test <- c(rep(1,N2),rep(2,N3),rep(3,N8))
# y.train <- cbind(g==1,g==2,g==3)*1
# y.test <- cbind(g.test==1,g.test==2,g.test==3)*1
# check all the row sums are 1
# all(apply(y.train,1,sum)==1)
lm1 <- lm(y.train[,1]~.,data=data.frame(x.train.fda))
lm2 <- lm(y.train[,2]~.,data=data.frame(x.train.fda))
lm3 <- lm(y.train[,3]~.,data=data.frame(x.train.fda))

# prediction
y.train.hat <- cbind(fitted(lm1),fitted(lm2),fitted(lm3))
g.hat <- apply(y.train.hat,1,which.max)
# train error
ols.train.error <- mean(g.hat!=g)

y.test.hat <- cbind(predict(lm1,data.frame(x.test.fda)),predict(lm2,dat
a.frame(x.test.fda)),predict(lm3,data.frame(x.test.fda)))
g.test.hat <- apply(y.test.hat,1,which.max)
# test error
ols.test.error <- mean(g.test.hat!=g.test)

# Logistic regression
ans.logistic <- glmnet(x.train.fda,g, family=c("multinomial"),lambda=0,
standardize=F)
g.hat.logist <- predict(ans.logistic,x.train.fda,type = "class")
# apply(predict(ans.logistic,x.test.fda,type = "response"),1,which.max)
g.test.hat.logist <- predict(ans.logistic,x.test.fda,type = "class")

# train error
logistic.train.error <- mean(g.hat.logist!=g)
# test error
logistic.test.error <- mean(g.test.hat.logist!=g.test)

# LDA
pi1.hat <- mean(g==1)
pi2.hat <- mean(g==2)
pi3.hat <- mean(g==3)

mu.1.hat <- apply(x.train.fda[1:n2,],2,mean)
mu.2.hat <- apply(x.train.fda[(n2+1):(n2+n3)],2,mean)
```

```

mu.3.hat <- apply(x.train.fda[(n2+n3+1):n,],2,mean)

S.w <- (t(x.train.fda[1:n2,] - rep(1,n2)%*% t(mu.1.hat)) %*% (x.train.fda[1:n2,] - rep(1,n2)%*% t(mu.1.hat)) +
      t(x.train.fda[(n2+1):(n2+n3),] - rep(1,n3)%*% t(mu.2.hat)) %*% (x.train.fda[(n2+1):(n2+n3),] - rep(1,n3)%*% t(mu.2.hat)) +
      t(x.train.fda[(n2+n3+1):n,] - rep(1,n8)%*% t(mu.3.hat)) %*% (x.train.fda[(n2+n3+1):n,] - rep(1,n8)%*% t(mu.3.hat)))/(n-3)

slope1 <- solve(S.w) %*% (mu.2.hat-mu.1.hat)
intercept1 <- log(pi2.hat/pi1.hat) - t(mu.2.hat+mu.1.hat) %*% solve(S.w) %*% (mu.2.hat-mu.1.hat)/2

slope2 <- solve(S.w) %*% (mu.3.hat-mu.1.hat)
intercept2 <- log(pi3.hat/pi1.hat) - t(mu.3.hat+mu.1.hat) %*% solve(S.w) %*% (mu.3.hat-mu.1.hat)/2

slope3 <- solve(S.w) %*% (mu.3.hat-mu.2.hat)
intercept3 <- log(pi3.hat/pi2.hat) - t(mu.3.hat+mu.2.hat) %*% solve(S.w) %*% (mu.3.hat-mu.2.hat)/2
# train error
delta.hat.train <-
cbind(log(pi1.hat)-diag((x.train.fda-rep(1,n)%*% t(mu.1.hat))%*% solve(S.w)%*%t(x.train.fda-rep(1,n)%*% t(mu.1.hat))/2),
      log(pi2.hat)-diag((x.train.fda-rep(1,n)%*% t(mu.2.hat))%*% solve(S.w)%*%t(x.train.fda-rep(1,n)%*% t(mu.2.hat))/2),
      log(pi3.hat)-diag((x.train.fda-rep(1,n)%*% t(mu.3.hat))%*% solve(S.w)%*%t(x.train.fda-rep(1,n)%*% t(mu.3.hat))/2))
g.hat.lda <- apply(delta.hat.train,1,which.max)

lda.train.error <- mean(g.hat.lda!=g)

# test error
N <- N2+N3+N8
delta.hat.test <-
cbind(log(pi1.hat)-diag((x.test.fda-rep(1,N)%*% t(mu.1.hat))%*% solve(S.w)%*%t(x.test.fda-rep(1,N)%*% t(mu.1.hat))/2),
      log(pi2.hat)-diag((x.test.fda-rep(1,N)%*% t(mu.2.hat))%*% solve(S.w)%*%t(x.test.fda-rep(1,N)%*% t(mu.2.hat))/2),
      log(pi3.hat)-diag((x.test.fda-rep(1,N)%*% t(mu.3.hat))%*% solve(S.w)%*%t(x.test.fda-rep(1,N)%*% t(mu.3.hat))/2))
g.test.hat.lda <- apply(delta.hat.test,1,which.max)
lda.test.error <- mean(g.test.hat.lda!=g.test)

# error table
error.table3 <- data.frame("OLS_FDA"=c(ols.train.error,ols.test.error),
"Logistic_FDA"=c(logistic.train.error,logistic.test.error),"LDA_FDA"=c(lda.train.error,lda.test.error))

```




```

row.names(error.table3)<-c("train error", 'test error')

# plot the first two PC scores and the decision boundaries
# OLS
x <- cbind(1,x.train.fda)
B <- solve(t(x)%*%x)%*%t(x)%*%y.train
par(mfrow=c(1,3))
plot(fst.score.fda,snd.score.fda,col=rep(2:4,c(n2,n3,n8)),pch=rep(2:4,c(
n2,n3,n8)),xlab="First Score",ylab="Second Score",main="OLS")
abline(b=(B[2,2]-B[2,1])/(B[3,1]-B[3,2]),a=(B[1,2]-B[1,1])/(B[3,1]-B[3,
2]))
abline(b=(B[2,3]-B[2,2])/(B[3,2]-B[3,3]),a=(B[1,3]-B[1,2])/(B[3,2]-B[3,
3]))
abline(b=(B[2,3]-B[2,1])/(B[3,1]-B[3,3]),a=(B[1,3]-B[1,1])/(B[3,1]-B[3,
3]))
# logistic
coef_logit <- coef(ans.logistic)
coef.log <- cbind(coef(ans.logistic)[[1]],coef(ans.logistic)[[2]],coef(
ans.logistic)[[3]])
plot(fst.score.fda,snd.score.fda,col=rep(2:4,c(n2,n3,n8)),pch=rep(2:4,c(
n2,n3,n8))
, xlab="First Score",ylab="Second Score",main="Logistic Regression"
)
abline(b=(coef.log[2,2]-coef.log[2,1])/(coef.log[3,1]-coef.log[3,2])
,a=(coef.log[1,2]-coef.log[1,1])/(coef.log[3,1]-coef.log[3,2]))
abline(b=(coef.log[2,3]-coef.log[2,2])/(coef.log[3,2]-coef.log[3,3])
,a=(coef.log[1,3]-coef.log[1,2])/(coef.log[3,2]-coef.log[3,3]))
abline(b=(coef.log[2,3]-coef.log[2,1])/(coef.log[3,1]-coef.log[3,3])
,a=(coef.log[1,3]-coef.log[1,1])/(coef.log[3,1]-coef.log[3,3]))
# LDA
plot(fst.score.fda,snd.score.fda,col=rep(2:4,c(n2,n3,n8)),pch=rep(2:4,c(
n2,n3,n8))
, xlab="First Score",ylab="Second Score",main="LDA")
abline(b=-slope1[1]/slope1[2],a=-intercept1/slope1[2])
abline(b=-slope2[1]/slope2[2],a=-intercept2/slope2[2])
abline(b=-slope3[1]/slope3[2],a=-intercept3/slope3[2])

```




```
par(mfrow=c(1,1))
error.table3

##               OLS_FDA Logistic_FDA    LDA_FDA
## train error 0.02278612  0.02123252 0.02382185
## test error  0.08490566  0.08490566 0.08867925
```

See the table.

(g)

```
error.table;error.table2;error.table3

##               OLS    Logistic    LDA
## train error 0.02278612 0.00000000 0.02382185
## test error  0.08490566 0.08301887 0.08867925

##               OLS_PCA Logistic_PCA    LDA_PCA
## train error 0.1532885    0.1460383 0.1589850
## test error  0.2396226    0.2377358 0.2320755

##               OLS_FDA Logistic_FDA    LDA_FDA
## train error 0.02278612  0.02123252 0.02382185
## test error  0.08490566  0.08490566 0.08867925
```

It is obvious that if using FDA to reduce the dimension, errors would not change except logistic regression. However, when using PCA, errors would increase significantly. Hence, for classification, FDA could beat PCA.

(h)

```
n.train <- n
set.seed(1)
nfolds <- 5
s <- split(sample(n.train),rep(1:nfolds,length=n.train))
p <- dim(x.svd$v)[1]
# Matrix to store predictions
y.cv <- matrix(NA, n.train,3)
g.cv <- rep(NA,n.train)

# CV error
cv.error.ols <- rep(NA,p)
for(i in 1:p){
  x.pc <- x.centered%%x.svd$v[,1:i]
  for (j in seq(nfolds)) {
    lmall <- lm(y.train[-s[[j]],]~.,data=data.frame(x.pc[-s[[j]],]))
    y.cv[s[[j]],] <- cbind(1,x.pc[s[[j]],]) %% coef(lmall)
  }
  g.cv <- apply(y.cv,1,which.max)
  cv.err <- mean(g.cv!=g)
  cv.error.ols[i] <- cv.err
}
```

```

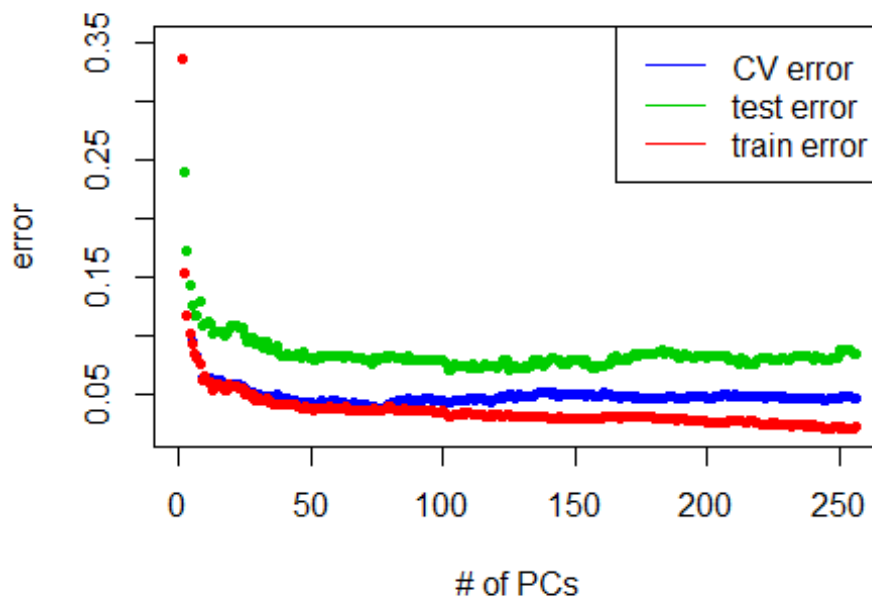
# test error
test.error.ols <- rep(NA,p)
for(i in 1:p){
  x.pc <- x.centered%%x.svd$v[,1:i]
  lmall <- lm(y.train~.,data=data.frame(x.pc))
  y.pred <- cbind(1,x.test.centered%%x.svd$v[,1:i]) %% coef(lmall)
  g.hat.test <- apply(y.pred,1,which.max)
  test.err <- mean(g.hat.test!=g.test)
  test.error.ols[i] <- test.err
}

# train error
train.error.ols <- rep(NA,p)
for(i in 1:p){
  x.pc <- x.centered%%x.svd$v[,1:i]
  lmall <- lm(y.train~.,data=data.frame(x.pc))
  y.pred <- cbind(1,x.test.centered%%x.svd$v[,1:i]) %% coef(lmall)
  g.hat.train <- apply(y.pred,1,which.max)
  train.err <- mean(g.hat.train!=g)
  train.error.ols[i] <- train.err
}

plot(seq(p),cv.error.ols,xlab="# of PCs", ylab="error",
     col=4,pch=20,ylim=c(0.02,0.35),main="Perform CV By Using OLS")
points(seq(p),test.error.ols,col=3,pch=20)
points(seq(p),train.error.ols,col=2,pch=20)
legend("topright",c("CV error","test error","train error"),col=c(4,3,2)
, lty=c(1,1))

```

Perform CV By Using OLS



```
which.min(cv.error.ols)
## [1] 76
# [1] 76
test.error.ols[76];train.error.ols[76]
## [1] 0.08113208
## [1] 0.03780425
which.min(test.error.ols)
## [1] 102
# [1] 102
test.error.ols[102];train.error.ols[102]
## [1] 0.07169811
## [1] 0.03262558
```

According to CV errors, we should keep 76 PC scores, and the training and test error are 0.03780425 and 0.08113208, respectively. (According to test errors, we should keep 102 PC scores, and the training and test error are 0.03262558 and 0.07169811)

(i)

```

# LDA
pi1.hat <- mean(g==1)
pi2.hat <- mean(g==2)
pi3.hat <- mean(g==3)
set.seed(1)
nfolds <- 5
s <- split(sample(n.train),rep(1:nfolds,length=n.train))
# CV error
cv.error.lda <- rep(NA,p)
delta.hat.cv <- matrix(NA,n,3)
for(i in 1:p){
  x.pc <- x.centered%%x.svd$v[,1:i]
  for (j in seq(nfolds)) {
    g1 <- x.pc[-s[[j]],,drop=F][g[-s[[j]]]==1,,drop=F]
    g2 <- x.pc[-s[[j]],,drop=F][g[-s[[j]]]==2,,drop=F]
    g3 <- x.pc[-s[[j]],,drop=F][g[-s[[j]]]==3,,drop=F]
    mu.1.hat <- apply(g1,2,mean)
    mu.2.hat <- apply(g2,2,mean)
    mu.3.hat <- apply(g3,2,mean)
    c1 <- sum(g[-s[[j]]]==1)
    c2 <- sum(g[-s[[j]]]==2)
    c3 <- sum(g[-s[[j]]]==3)
    S.w <- (t(g1 - rep(1,c1))%*% t(mu.1.hat))%*%(g1 - rep(1,c1))%*% t(mu.
1.hat)) +
      t(g2 - rep(1,c2))%*% t(mu.2.hat))%*%(g2 - rep(1,c2))%*% t(m
u.2.hat)) +
      t(g3 - rep(1,c3))%*% t(mu.3.hat))%*%(g3 - rep(1,c3))%*%
t(mu.3.hat)))/(c1+c2+c3-3)

    tt <- x.pc[s[[j]],,drop=F]; nn <- length(s[[j]])
    delta.hat <-
      cbind(log(pi1.hat)-diag((tt-rep(1,nn))%*% t(mu.1.hat))%*%
        solve(S.w)%*%t(tt-rep(1,nn))%*%t(mu.1.hat)
)/2)
      ,log(pi2.hat)-diag((tt-rep(1,nn))%*% t(mu.2.hat))%*%
        solve(S.w)%*%t(tt-rep(1,nn))%*%t(mu.2.hat
))/2)
      ,log(pi3.hat)-diag((tt-rep(1,nn))%*% t(mu.3.hat))%*%
        solve(S.w)%*%t(tt-rep(1,nn))%*%t(mu.3.hat
))/2))
    delta.hat.cv[s[[j]],] <- delta.hat
  }
  g.lda.hat <- apply(delta.hat.cv,1,which.max)
  cv.err <- mean(g.lda.hat!=g)
  cv.error.lda[i] <- cv.err
}

# train and test error
train.error.lda <- rep(NA,p)
test.error.lda <- rep(NA,p)

```


```

for(i in 1:p){
  x.pc <- x.centered%%x.svd$v[,1:i]
  x.pc.test <- x.test.centered%%x.svd$v[,1:i]
  g1 <- x.pc[g==1,,drop=F]
  g2 <- x.pc[g==2,,drop=F]
  g3 <- x.pc[g==3,,drop=F]
  mu.1.hat <- apply(g1,2,mean)
  mu.2.hat <- apply(g2,2,mean)
  mu.3.hat <- apply(g3,2,mean)
  S.w <- (t(x.pc[1:n2,] - rep(1,n2)%% t(mu.1.hat))
    %% (x.pc[1:n2,] - rep(1,n2)%% t(mu.1.hat))
    + t(x.pc[(n2+1):(n2+n3),] - rep(1,n3)%% t(mu.2.hat)) %%
    (x.pc[(n2+1):(n2+n3),] - rep(1,n3)%% t(mu.2.hat)) +
    t(x.pc[(n2+n3+1):n,] - rep(1,n8)%% t(mu.3.hat))
    %% (x.pc[(n2+n3+1):n,] - rep(1,n8)%% t(mu.3.hat)))/(n-3)
  delta.hat.train <-
    cbind(log(pi1.hat)-diag((x.pc-rep(1,n)%% t(mu.1.hat))%%
      solve(S.w)%%t(x.pc-rep(1,n)%% t(mu.1.ha
t))/2),
      log(pi2.hat)-diag((x.pc-rep(1,n)%% t(mu.2.hat))%%
      solve(S.w)%%t(x.pc-rep(1,n)%% t(mu.2.ha
t))/2),
      log(pi3.hat)-diag((x.pc-rep(1,n)%% t(mu.3.hat))%%
      solve(S.w)%%t(x.pc-rep(1,n)%% t(mu.3.ha
t))/2))
  delta.hat.test <-
    cbind(log(pi1.hat)-diag((x.pc.test-rep(1,N)%% t(mu.1.hat))%%
      solve(S.w)%%t(x.pc.test-rep(1,N)%% t(mu
.1.hat))/2),
      log(pi2.hat)-diag((x.pc.test-rep(1,N)%% t(mu.2.hat))%%
      solve(S.w)%%t(x.pc.test-rep(1,N)%% t(mu
.2.hat))/2),
      log(pi3.hat)-diag((x.pc.test-rep(1,N)%% t(mu.3.hat))%%
      solve(S.w)%%t(x.pc.test-rep(1,N)%% t(mu
.3.hat))/2))
  g.hat.train <- apply(delta.hat.train,1,which.max)
  g.hat.test <- apply(delta.hat.test,1,which.max)
  train.err <- mean(g.hat.train!=g)
  test.err <- mean(g.hat.test!=g.test)
  train.error.lda[i] <- train.err ; test.error.lda[i] <- test.err
}

plot(seq(p),cv.error.lda,xlab="# of PCs", ylab="error",
      col=4,pch=20,ylim=c(0.02,0.35),main="Perform CV By Using LDA")
points(seq(p),test.error.lda,col=3,pch=20)
points(seq(p),train.error.lda,col=2,pch=20)
legend("topright",c("CV error","test error","train error"),col=c(4,3,2)
,lty=c(1,1))

```

Perform CV By Using LDA



```
which.min(cv.error.lda)
## [1] 78
#[1] 78
test.error.lda[78];train.error.lda[78]
## [1] 0.08490566
## [1] 0.03728638
which.min(test.error.lda)
## [1] 126
#[1] 126
test.error.lda[126];train.error.lda[126]
## [1] 0.07169811
## [1] 0.03469705
```

According to CV errors, we should keep 78 PC scores, and the training and test error are 0.03728638 and 0.08490566, respectively. (According to test errors, we should keep 126 PC scores, and the training and test error are 0.03469705 and 0.07169811)

(j)

```

set.seed(1)
nfolds <- 5
s <- split(sample(n.train),rep(1:nfolds,length=n.train))
# CV error
cv.error.log <- rep(NA,p)
g.hat.log <- rep(NA,n)
for(i in 2:p){
  x.pc <- x.centered%%x.svd$v[,1:i]
  for (j in seq(nfolds)) {
    ans.logistic <- glmnet(x.pc[-s[[j]],,drop=F],g[-s[[j]]], family=c("
multinomial")
                                ,lambda=0,standardize=F)
    g.hat.log[s[[j]]] <- predict(ans.logistic,x.pc[s[[j]],,drop=F],type
= "class")
  }
  cv.err <- mean(g.hat.log!=g)
  cv.error.log[i] <- cv.err
}

# train error and test error
train.error.log <- rep(NA,p)
test.error.log <- rep(NA,p)

for(i in 2:p){
  x.pc <- x.centered%%x.svd$v[,1:i]
  x.pc.test <- x.test.centered%%x.svd$v[,1:i]
  ans.logistic <- glmnet(x.pc,g, family=c("multinomial"),lambda=0,stand
ardize=F)
  g.hat.log <- predict(ans.logistic,x.pc,type = "class")
  g.hat.test.log <- predict(ans.logistic,x.pc.test,type = "class")
  train.err <- mean(g.hat.log!=g)
  test.err <- mean(g.hat.test.log!=g.test)
  train.error.log[i] <- cv.err
  test.error.log[i] <- test.err
}

```

(k)

```

test.error.ols[76];train.error.ols[76]
## [1] 0.08113208
## [1] 0.03780425
test.error.ols[102];train.error.ols[102]
## [1] 0.07169811
## [1] 0.03262558
test.error.lda[78];train.error.lda[78]

```

```
## [1] 0.08490566
## [1] 0.03728638
test.error.lda[126];train.error.lda[126]
## [1] 0.07169811
## [1] 0.03469705
```

Based on the results we obtained, if using FDA to reduce the dimension, we simply need the first two FDA scores to achieve comparable test errors of minimum test errors obtained by using PC scores. Therefore, FDA is better for classification.